

## 🔧 Troubleshooting Kubernetes Performance Issues in Production: A Comprehensive Step-by-Step Guide 🔧

Kubernetes has revolutionized how we manage containerized applications, but when performance issues arise in production, it can be challenging to diagnose and resolve them quickly. Having navigated these waters, I'd like to share an expanded step-by-step guide that can help you debug Kubernetes issues effectively.



### Troubleshooting **KUBERNETES** PRODUCTION



1. **Service Down:** Are any services completely down, experiencing sluggishness, or intermittently failing? Gather user feedback to assess impact.
2. **Pod Issues:** Are there problems with pod startup times, scheduling delays, or unexpected crashes? Understanding the context of these issues can provide clues to the root cause.
3. **Network Problems:** Are you facing any challenges with inter-pod communication, external access, or DNS resolution? A broken network can often lead to cascading failures across services.
4. **Resource Exhaustion:** Is the system reaching its limits for CPU, memory, or disk space? Look for signs of throttling or high utilization that could affect application performance.
5. **Configuration Changes:** Have there been any recent configuration changes to your Kubernetes cluster, deployments, or services? Understanding the timeline of changes can help you pinpoint new issues.
6. **External Dependencies:** Are there any third-party services or APIs that your applications depend on? If these are down or slow, they can significantly impact your performance.
7. **User Traffic Patterns:** Has there been a spike in user traffic? An unexpected increase in load can expose underlying performance issues that were previously unnoticed.
8. **Historical Trends:** Are there historical trends in application performance that correlate with the current issue? Reviewing past incidents may provide insights into potential recurring problems.

### Step 1: Monitor Resource Utilization

Start by checking the resource utilization of your pods and nodes. Running low on resources can lead to chaos. To assess our resource usage, I run:

```
kubectl top pods --all-namespaces
```

```
kubectl top nodes
```

Look for pods that are nearing their resource limits, as this can lead to throttling and degraded performance.

### Step 2: Analyze Logs

Next, dive into the logs of the affected pods. You can use the following command to access the logs:

```
kubectl logs <pod-name> -n <namespace>
```

Look for error messages, warnings, or unusual patterns that may indicate what's going wrong.

**1.If the pod is failing due to application errors or crashes, check its logs:**

```
kubectl logs <pod_name> -n <namespace>
```

**2. If the pod has multiple containers, specify the container name:**

```
kubectl logs <pod_name> -c <container_name> -n <namespace>
```

**3. For pods in a crash loop, use the --previous flag to see logs from the previous instance:**

```
kubectl logs <pod_name> --previous -n <namespace>
```

### Step 3: Check Events

Kubernetes events can provide insights into what's happening within your cluster. Run the following command to view recent events:

```
kubectl get events --namespace=<namespace>
```

This can help identify issues such as failed scheduling or resource contention.

Kubernetes events often provide valuable context for what's going wrong:

```
kubectl get events -n <namespace> --sort-by='.lastTimestamp'
```

### Step 4: Inspect Resource Requests and Limits

Ensure that your resource requests and limits are properly set. If they are too low, it can lead to performance issues; if they are too high, they can cause resource contention.

Review your deployment YAML files, and adjust requests and limits as necessary:

```
resources:
```

```
  requests:
```

```
    cpu: "250m"
```

```
    memory: "256Mi"
```

```
  limits:
```

```
    cpu: "750m"
```

```
    memory: "512Mi"
```

### Step 5: Analyze Network Performance

Network issues can also affect performance. Check for network latency and connectivity problems by using tools like ping or curl between pods.

You can also use `kubectl exec` to run these commands directly within the container:

```
kubectl exec -it <pod-name> -- curl http://<other-pod-ip>:<port>
```

### Step 6: Review Pod Health

Make sure your pods are healthy. Check the status of your pods to ensure they are not in a crash loop or experiencing restarts:

```
kubectl get pods -n <namespace>
```

Use `kubectl describe pod <pod-name>` to see detailed information about the pod's state and any error messages.

**Look for pod statuses like CrashLoopBackOff, Error, Pending, OOMKilled, or Completed.**

### Step 7: Investigate Node Health

When I notice multiple pods acting up, the first thing I do is check the health of the nodes:

```
kubectl get nodes
```

```
kubectl describe node <node_name>
```

I always keep an eye out for DiskPressure or MemoryPressure alerts, as these are crucial indicators of underlying problems.

### Step 8: Network and DNS Issues

Next, I verify pod communication and check network policies:

```
kubectl describe networkpolicy -n <namespace>
```

For DNS, I ensure the CoreDNS pods are operational with:

```
kubectl get pods -n kube-system | grep coredns
```

```
kubectl logs <coredns_pod_name> -n kube-system
```

### Step 9: Service and Endpoint Issues

I check the status of our services to identify any potential problems:

```
kubectl get svc -n <namespace>
```

If something seems off, I dig deeper:

```
kubectl describe svc <service_name> -n <namespace>
```

### Step 10: Investigate Ingress and Load Balancer

Next, I confirm that our Ingress and LoadBalancer configurations are set up correctly:

```
kubectl get ingress -n <namespace>
```

```
kubectl describe ingress <ingress_name> -n <namespace>
```

```
kubectl logs <ingress-controller-pod> -n <namespace>
```

### Step 11: Persistent Volume and Storage Issues

If I encounter stuck pods, I immediately check the PVC and PV statuses:

```
kubectl describe pvc <pvc_name> -n <namespace>
```

### Step 12: Inspect Deployment and StatefulSet Rollouts

Keeping track of application deployments is crucial. I check the rollout status like this:

```
kubectl rollout status deployment <deployment_name> -n <namespace>
```

### Step 13: Resource Quotas and Limits

If scheduling issues arise, I check the resource quotas for the namespace:

```
kubectl get resourcequotas -n <namespace>
```

```
kubectl describe resourcequota <quota_name> -n <namespace>
```